

# Lecture 3: Data Structures in R



UNIVERSITY OF  
SAN FRANCISCO

James D. Wilson  
ICPSR: Network Analysis I



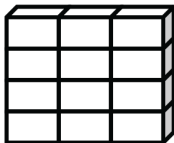
- A **data structure** is a format or organization of data in software that enables efficient use.
- Every programming language has its own types of data structures
- In  $\mathbb{R}$ , you can create your own type of data structure; however, there are some that are automatically recognized by the software.
- **Examples:** list, array, data.frame, vector, matrix, string




(a) Vector



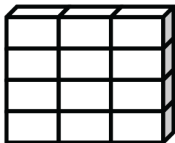
(b) Matrix



(c) Array

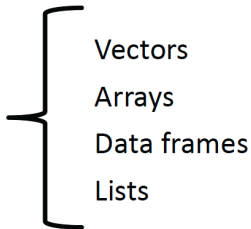


(d) Data frame



Columns can be different modes

(e) List





Dimension	Homogeneous	Heterogeneous
1	Atomic Vector	List
2	Matrix	Data Frame
$n$	Array	

**Homogeneous:** All contents must be of the same type

**Heterogeneous:** Contents can be of different types

**Note:** There are no 0-dimensional (scalar) types in R, only vectors of length one

# Part I: Vectors



- The basic data structure in R is the vector
- There two types of vectors: atomic vectors and lists

## Properties of Vectors

- Type (`typeof()`)
- Length (`length()`)
- Attributes (`attributes()`)

Use `is.atomic()` or `is.list()` to determine if an object is a vector, **not** `is.vector()`



## Four Common Types of Vectors

- Logical
- Integer
- Double (numeric)
- Character

```
> doubleAtomicVector <- c(1, 3.14, 99.999)
```

```
# use L prefix to get integers instead of doubles
```

```
> integerAtomicVector <- c(1L, 3L, 19L)
```

```
> logicalAtomicVector <- c(TRUE, FALSE, T, F)
```

```
> characterAtomicVector <- c("this", "is a", "string")
```

# Example: Try This



- 1 Create the vector `myFavNum` of you favorite fractional number
- 2 Create the vector `myNums` of your seven favorite numbers
- 3 Create the vector `firstNames` of the first names of two people next to you
- 4 Create the vector `myVec` of the last name and age of someone you know



# Example: Answer these



- 1 Guess and then check what types your vectors are.
- 2 Check the length of each vector.
- 3 Did you write the code in the console window or the editor?
- 4 How do you execute a line of code in the editor?
- 5 How do you execute multiple lines of code simultaneously in the editor?
- 6 Did you leverage the `TAB` button for auto-completion?

# Accessing Elements of a Vector



- To access the individual elements of a vector

```
> (myAtomicVector <- c(1, 2, 3, 4, -99, 5, NA, 4, 22.223))  
[1] 1.000 2.000 3.000 4.000 -99.000 5.000 NA  
[8] 4.000 22.223
```

#look at fifth element of the vector

```
> myAtomicVector[5]  
[1] -99
```

```
> myAtomicVector[c(1, 2, 5, 9)]  
[1] 1.000 2.000 -99.000 22.223
```

```
> myAtomicVector[10]  
[1] NA
```

#look at the third through eighth elements of the vector

```
> myAtomicVector[3:8]  
[1] 3 4 -99 5 NA 4
```



- To look at the first and last 6 elements of a vector

```
> (myAtomicVector <- c(1, 2, 3, 4, -99, 5, NA, 4, 22.223))  
[1] 1.000 2.000 3.000 4.000 -99.000 5.000 NA  
[8] 4.000 22.223
```

#look at the first and last six elements of the vector

```
> head(myAtomicVector)  
[1] 1.000 2.000 3.000 4.000 -99.000 5.000
```

```
> tail(myAtomicVector)  
[1] 4.000 -99.000 5.000 NA 4.000 22.223
```



- ➊ Add `myFavNum` to the seventh entry of `myNums` and store the result in a variable named `myFirstAddition`
- ➋ Add `myFavNum` to each of the seven entries of `myNums` and store the result in a variable named `mySecondAddition`
- ➌ Add `myFavNum` to **all** of the values in `myNums` and store the result in a variable named `myFirstSum`
- ➍ Add `myFavNum` to the smallest number in `myNums` and store the result in a variable named `thisIsGettingMoreComplex`
- ➎ Add the second entry of `myNums` to the age of the person you select for `myVec` and store the result in a variable named `whatTypeOfVectorIsThis`
  - Does what we did make sense? Did it work? Why?



```
# preamble
myFavNum <- 3.1415
myNums <- c(1, 3, 55, 33, 86, -sqrt(2), -110)
# also works myNums <- 1:7
firstNames <- c("Jeff", "Terence", "David")
myVec <- c("Parr", 99)
```

- ❶ myFirstAddition <- myFavNum + myNums[7]
- ❷ mySecondAddition <- myFavNum + myNums
- ❸ myFirstSum <- myFavNum + sum(myNums)
- ❹ thisIsGettingMoreComplex <- myFavNum + min(myNums)
- ❺ whatTypeOfVectorIsThis <- sum(c(myNums[2], myVec[2]))  
Error in sum(c(myNums[2], myVec[2])) :  
invalid 'type' (character) of argument



Missing values are specified with `NA`, a logical vector of length one.

- `NA` will always be coerced to the correct type if used inside `c()`

```
> c(1, 2, 3, NA)
[1] 1 2 3 NA
```

```
> x[1]
[1] 1
```

```
> x <- c(1, 2, 3, NA)
```

```
> x[4]
[1] NA
```

```
> typeof(x)
[1] "double"
```

```
> typeof(x[4])
[1] "double"
```



- Certain functions will fail when applied to vectors with an `NA`

```
> myAtomicVector_01 <- c(99.1, 98.2, 97.3, 96.4, NA)
[1] 99.1 98.2 97.3 96.4 NA
```

```
> sum(myAtomicVector_01)
[1] NA
```

```
> mean(myAtomicVector_01)
[1] NA
```



- You can avoid this by providing the argument `na.rm = TRUE`

```
> sum(myAtomicVector_01, na.rm = TRUE)
[1] 391
```

```
> mean(myAtomicVector_01, na.rm = TRUE)
[1] 97.75
```





To check the type of a vector, use `typeof()`, or more specifically

- `is.character()`
- `is.double()`
- `is.integer()`
- `is.logical()`
- `is.na()`



Coercion is a great feature in  $\mathbb{R}$  which can make coding easy, but may also have unintended consequences.

- All elements in an atomic vector must be the same type
- If you attempt to combine different types in an atomic vector they will be coerced to the most flexible type
- **Most to least flexible types** ↓
  - character
  - double
  - integer
  - logical

- When a logical vector is coerced to numeric (double or integer), TRUE = 1 and FALSE = 0

```
> x <- c("abc", 123)
```

```
> typeof(x)
```

```
[1] "character"
```

You can explicitly coerce using `as.character()`, `as.double()`, `as.integer()`, and `as.logical()`



- A quick way to figure out what data structure an object is composed of is to use `str()`, which is short for structure
- `str()` provides a concise description for any R data structure



- The syntax is awkward and takes some time to get used to
- Once you understand the sequence of events in conditional subsetting, it will feel more natural
- Try to figure out what is happening in the following example:

```
> (myAtomicVector_01 <- c(99.1, 98.2, 97.3, 96.4))  
[1] 99.1 98.2 97.3 96.4
```

```
> myAtomicVector_01[myAtomicVector_01 > 98]  
[1] 99.1 98.2
```



What is actually happening in the last slide:

- 1 The `myAtomicVector_01 > 98` part of the statement tests each element of the vector to see whether it is  $> 98$  and returns a `LOGICAL` value for each test which, in this case, returns the logical vector `(T T F F)`
- 2 The vector `(T T F F)` is passed to `myAtomicVector_01`, which returns the first two elements and omits the final two
  - An equivalent statement would be  
`myAtomicVector_01[c(T, T, F, F)]`



---

Function	Action
----------	--------

---

<code>seq(from, to, by)</code>	Creates a vector of numbers from <code>from</code> to <code>to</code> in increments of <code>by</code>
--------------------------------	--

<code>rep(x, times)</code>	Creates a vector that repeats the values in <code>x</code> exactly <code>times</code> number of times
----------------------------	---

<code>x + (-, /, *) y</code>	For <code>x</code> and <code>y</code> of the <i>same length</i> , calculates a vector of the same length where each entry is the <b>entry-wise</b> summation (subtraction, division, or product) of <code>x</code> and <code>y</code>
------------------------------	---

---



- If you would like to create a vector that is a sequence of numbers from  $x$  to  $y$  that increase by exactly one, then you can simply write

$x:y$

- `rep()` can be applied to a `seq()`, providing a flexible means to create sequences with repeating patterns.

## Example:

```
> rep(seq(1, 1.3, .1), 2)
[1] 1.0 1.1 1.2 1.3 1.0 1.1 1.2 1.3
```



# Example



```
> x <- rep(c(1,2), 3)
```

```
> y <- seq(from = .5, to = 3, by = .5)
```

```
> x
```

```
[1] 1 2 1 2 1 2
```

```
> y
```

```
[1] 0.5 1.0 1.5 2.0 2.5 3.0
```

```
> x+y
```

```
[1] 1.5 3.0 2.5 4.0 3.5 5.0
```

```
> x/y
```

```
[1] 2.0000000 2.0000000 0.6666667 1.0000000 0.4000000 0.6666667
```

# A List of Logical Operators



Operator	Description
<code>&lt;</code>	Less than
<code>&lt;=</code>	Less than or equal to
<code>&gt;</code>	Greater than
<code>&gt;=</code>	Greater than or equal to
<code>==</code>	Exactly equal to
<code>!=</code>	Not equal to
<code>!x</code>	Not $x$
<code>x   y</code>	$x$ or $y$
<code>x &amp; y</code>	$x$ and $y$
<code>isTRUE(x)</code>	Test if $x$ is TRUE



- A name is a vector **attribute**
- Can be identified using the `names()` function

```
> x <- c(1, 2, 3)
> names(x)
NULL
```

```
> x <- c(1, 2, 3); names(x) <- c("a", "b", "c")
> names(x)
[1] "a" "b" "c"
```

```
> x <- c(a = 1, b = 2, c = 3)
> names(x)
[1] "a" "b" "c"
```

```
> x <- c(a = 1, b = 2, 3)
> names(x)
[1] "a" "b" ""
```

## Part II: Matrices and Arrays



- By giving an atomic vector a dimension attribute, it behaves like a multi-dimensional array
- A special case of the array is a matrix, a two-dimensional array
- A matrix has 2 dimensions, and an array has  $n \geq 2$  - dimensions.
- Matrices and arrays are created with `matrix()` and `array()`



```
> x <- matrix(1:10, nrow = 2, ncol = 5)
# can drop nrow and ncol to shorten but keep in this order
```

```
> x
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	1	3	5	7	9
[2,]	2	4	6	8	10



```
> y <- array(1:12, c(2, 3, 2))
```

```
> y
```

```
, , 1
```

```
      [,1] [,2] [,3]
```

```
[1,]     1     3     5
```

```
[2,]     2     4     6
```

```
, , 2
```

```
      [,1] [,2] [,3]
```

```
[1,]     7     9    11
```

```
[2,]     8    10    12
```



---

## 1-D Function    n-D Functions

---

<code>length()</code>	<code>nrow()</code> , <code>ncol()</code> , <code>dim()</code>
<code>names()</code>	<code>rownames()</code> , <code>colnames()</code> , <code>dimnames()</code>
<code>c()</code>	<code>cbind()</code> , <code>rbind()</code>

---

**Note:** a matrix or array can also be one-dimensional, e.g., an object that is defined as a matrix is permitted to only have one column or one row; although they may look and behave alike, a vector and a one-dimensional matrix behave differently and may generate strange output when using certain functions, e.g., `tapply()`



# Common R Functions for Working with Data



Function	Purpose
<code>length(object)</code>	Number of elements/components.
<code>dim(object)</code>	Dimensions of an object.
<code>str(object)</code>	Structure of an object.
<code>class(object)</code>	Class or type of an object.
<code>mode(object)</code>	How an object is stored.
<code>names(object)</code>	Names of components in an object.
<code>c(object, object, ...)</code>	Combines objects into a vector.
<code>cbind(object, object, ...)</code>	Combines objects as columns.
<code>rbind(object, object, ...)</code>	Combines objects as rows.
<code>object</code>	Prints the object.
<code>head(object)</code>	Lists the first part of the object.
<code>tail(object)</code>	Lists the last part of the object.
<code>ls()</code>	Lists current objects.
<code>rm(object, object, ...)</code>	Deletes one or more objects. The statement <code>rm(list = ls())</code> will remove most objects from the working environment.
<code>newobject &lt;- edit(object)</code>	Edits object and saves as newobject.
<code>fix(object)</code>	Edits in place.

# Part III: Lists



- Lists are different from atomic vectors as elements of a list can be of any type, including lists
- A list is constructed using `list()` instead of `c()`

```
> myList <- list(10:12, "abc", c(3.1415, 9), c(T, F, F, F))

> str(myList)
List of 4
 $ : int [1:3] 10 11 12
 $ : chr "abc"
 $ : num [1:2] 3.14 9
 $ : logi [1:4] TRUE FALSE FALSE FALSE
```



- Lists are recursive, i.e., a list can contain lists, making them fundamentally different from atomic vectors
- Handy functions

---

Function	Action
<code>is.list()</code>	test if list
<code>as.list()</code>	coerce to list
<code>unlist()</code>	convert to atomic vector + coercion

---



- Entries in a list can contain any type of data structure
- To call a single entry (say the second one) in the list `myList`, use double brackets: `myList[[2]]`
- To call multiple entries in a list (say the first and second), use single brackets: `myList[1:2]`
- If the entries in a list are named, you can call them directly using `myList$Name`

# Subsetting Example



```
> myList <- list(10:12, Letters = "abc", c(3.1415, 9), Logicals =  
c(T, F, F, F))
```

```
> myList[[2]]  
[1] "abc"
```

```
> myList$Logicals  
[1] TRUE FALSE FALSE FALSE
```

```
> myList[1:2]  
[[1]]  
[1] 10 11 12
```

```
$Letters  
[1] "abc"
```

# Part IV: Data Frames



- Most common way of storing data in R
- A data frame is a list with equal-length vectors
- Each vector must be of the same data type

This is why we use





# Data Frame Summary Example



Summary of Data `ToothGrowth`: a data frame with 60 observations on 3 variables.

- `[,1]` len numeric: Tooth length
- `[,2]` supp factor: Supplement type (VC or OJ)
- `[,3]` dose numeric: Dose in milligrams/day

```
> str(ToothGrowth)
'data.frame': 60 obs. of  3 variables:
 $ len : num  4.2 11.5 7.3 5.8 6.4 10 11.2 11.2 5.2 7 ...
 $ supp: Factor w/ 2 levels "OJ","VC": 2 2 2 2 2 2 2 2 2 2 ...
 $ dose: num  0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 ...

> ?ToothGrowth
```

# Creating Data Frames



Create a data frame using `data.frame()`

```
# this is sloppy coding etiquette and is only for exposition
```

```
> (xyz <- data.frame(1:3, c("a", "b", "c")))
```

```
  X1.3 c..a....b....c..
```

```
1      1                a
```

```
2      2                b
```

```
3      3                c
```

```
> str(xyz)
```

```
'data.frame': 3 obs. of  2 variables:
```

```
$ X1.3      : int  1 2 3
```

```
$ c..a....b....c...: Factor w/ 3 levels "a","b","c": 1 2 3
```



Create a data frame using `data.frame()`

- Surround code with `()` to automatically print the result to the console
- After creating the data frame, the first column of untitled numbers are row numbers
- Observe that even though the entries in `letterColumn` are characters that an `str(letterColumn)` shows the column to be a `Factor`



- If you want to suppress R's default behavior of turning strings into factors, use the options `stringsAsFactors = FALSE`

```
> (xyz <- data.frame(numberColumn = 1:3, letterColumn = c("a", "b", "c"),  
  stringsAsFactors = F))
```

```
  numberColumn letterColumn  
1             1           a  
2             2           b  
3             3           c
```

```
> str(xyz)
```

```
'data.frame': 3 obs. of  2 variables:  
 $ numberColumn: int  1 2 3  
 $ letterColumn: chr  "a" "b" "c"
```



- **Note:** A data frame is a list, which means that `typeof(myDataFrame)` will output a list
- Instead use `class()` or `is.data.frame()`
- An object can be coerced to a data frame using `as.data.frame()`



- **When a data frame already exists**, you can easily combine/append another data frame or a vector to the original data frame
  - 1 Use `cbind()` to column-bind two data frames
    - **Note:** the number of columns in each data frame must be equal, and row names are ignored
  - 2 Use `rbind()` to row-bind two data frames
    - **Note:** the **number** and the **names** of columns must match

# Examples: `cbind()`



```
> (myDataFrame_01 <- data.frame(x = 1:3, y = c("A", "B", "c")))  
  x y  
1 1 A  
2 2 B  
3 3 c  
  
> (myDataFrame_02 <- cbind(myDataFrame_01, data.frame(z = -1:-3)))  
  x y  z  
1 1 A -1  
2 2 B -2  
3 3 c -3
```

# Examples: `rbind()`



```
> (myDataFrame_05 <- data.frame(x = 1:3, y = 98:100, z = 1000:1002))  
  x    y    z  
1 1   98 1000  
2 2   99 1001  
3 3  100 1002
```

```
> (myDataFrame_06 <- rbind(myDataFrame_05, qq = -1:-3))  
  x    y    z  
1  1   98 1000  
2  2   99 1001  
3  3  100 1002  
qq -1  -2   -3
```



# Example: Try these



```
> myDataFrame_05 <- data.frame(x = 1:3, y = 98:100, z = 1000:1002)

> myDataFrame_06 <- rbind(myDataFrame_05, ???)
```

- Based on the `myDataFrame_06` code, what happens if we replace `???` with:

- (a) `qqq = -1`
- (b) `qqq = -1:-2`
- (c) `qqq = -1:-99`
- (d) `qqq = c(-1, -2)`
- (e) `qqq = c("-1", -2)`
- (f) `qqq = c("a", -2, -3))`



```
> myDataFrame_05 <- data.frame(x = 1:3, y = 98:100, z = 1000:1002)
```

```
> myDataFrame_06 <- rbind(myDataFrame_05, ???)
```

- (a) Entire additional row of -1's
- (b) Entire additional row of repeating -1's and -2's
- (c) Additional row: -1, -2, -3
- (d) Entire additional row of repeating -1's and -2's
- (e) Entire additional row of repeating -1's and -2's as **characters** (non numeric), thereby changing **all** all data frame column types to **characters**
- (f) Additional row: a, -2, -3 as **characters** (non numeric), thereby changing **all** all data frame columns types to **characters**



- Use `cbind()` to column-bind a data frame with a vector
  - **Note:** This will only work if the vector has the same length as the number of rows in the data frame.

```
> (myDataFrame_07 <- data.frame(x = 1:3, y = 98:100, z = 1000:1002))
```

```
  x    y    z
1 1   98 1000
2 2   99 1001
3 3  100 1002
```

```
> (myDataFrame_08 <- cbind(myDataFrame_05, qq = -1:-3))
```

```
  x    y    z qq
1 1   98 1000 -1
2 2   99 1001 -2
3 3  100 1002 -3
```

# Example: Try these



```
> myDataFrame_07 <- data.frame(x = 1:3, y = 98:100, z = 1000:1002)

> myDataFrame_08 <- cbind(myDataFrame_07, ???)
```

- Based on the `myDataFrame_08` code, what happens if we replace `???` with:

- (a) `qqq = -1`
- (b) `qqq = -1:-2`
- (c) `qqq = -1:-99`
- (d) `qqq = c("-1", -2)`
- (e) `qqq = c("a", -2, -3))`



```
> myDataFrame_07 <- data.frame(x = 1:3, y = 98:100, z = 1000:1002)

> myDataFrame_08 <- cbind(myDataFrame_05, ???)
```

- (a) Entire additional column of `-1`'s
- (b) `<arguments imply differing number of rows: 3, 2>`
- (c) Extends the length of all other columns and repeats those values until `-99`
- (d) `<arguments imply differing number of rows: 3, 2>`
- (e) `<arguments imply differing number of rows: 3, 2>`
- (f) Additional column: `a`, `-2`, `-3` as **factors** (non numeric)